a **clear, grounded explanation** of what the **Agentic AI Lead interview** will actually be about — not in theory, but in **practice**, so you know what to expect and how to position yourself.

What this interview is *really* for
This interview is **not** about:

- Testing whether you know what agentic AI is
- Deep ML theory
- Coding ability
- Tool or framework trivia

You already passed the CTO. That means:

- They trust your seniority
- They trust your architectural mindset
- They believe you could add value

The Agentic AI Lead interview exists to answer **one specific question**:
"Can we trust Samuel to help us evolve agentic AI in a live, enterprise platform without introducing uncontrolled risk?"
Everything in the conversation maps back to that.

The Agentic AI Lead's mental checklist
Juber will be subconsciously assessing you across **five dimensions**.

---

1. Do you understand *agentic AI beyond the buzzword*?
He will want to see that you understand:

- Agentic AI is about **behaviour over time**
- Risk comes from **coordination and autonomy**, not just models
- The danger is not "wrong answers" but **unbounded behaviour**

He's listening for signals like:

- "orchestration"
- "decision boundaries"
- "human-in-the-loop"
- "observability"
- "guardrails"

If you talk about agentic AI as:

- "autonomous systems making decisions"
  that is a red flag.

---

2. Can you reason about control, not just intelligence?
This is the core of the interview.
He will probe:

- Where agents should act
- Where they must stop
- How behaviour is constrained
- How mistakes are contained

Expect questions like:

- "Where would you never allow an agent to act?"
- "What happens when an agent behaves unexpectedly?"
- "How do you roll back agent behaviour safely?"

He is testing whether you think in terms of:
blast radius, reversibility, and safety

---

3. Do you think in production terms, not design diagrams?
Juber is not interested in:

- Ideal architectures
- Clean-slate designs

He cares about:

- Live systems
- Customers already onboarded
- Changes with real consequences

That's why your **current role at SymphonyAI** matters.
He will listen for:

- Present-tense examples
- Trade-offs you made
- Decisions you *didn't* take
- What broke and how you fixed it

---

4. Can you align agentic AI with the Rezolve / Brain platform strategy?
This is where your CTO alignment matters.
He wants to see that you understand:

- Rezolve has a **shared Brain**
- Multiple products sit on top of it
- Agentic behaviour must not fragment that shared intelligence

So when you talk about agents, he's thinking:

- "Will this break reuse?"
- "Will this create bespoke logic?"
- "Can this be governed across customers?"

Your role here is to sound like someone who:
protects the platform while enabling progress

---

5. Do you show judgement and restraint?
This is subtle but critical.
Strong candidates talk about:

- What they enabled

Senior candidates also talk about:

- What they deliberately *didn't* enable
- When they slowed things down
- When they said no to autonomy

He is listening for:

- Wisdom
- Caution
- Long-term thinking

Not excitement.

What the conversation will *feel like*
It will feel:

- Exploratory
- Architectural
- Scenario-driven

It will **not** feel:

- Like a quiz
- Like an exam
- Like a sales pitch

He may:

- Interrupt you
- Ask "why?" multiple times
- Push on edge cases
- Challenge your assumptions

That's a good sign.

How you should mentally approach it
Do **not** go in thinking:
"I need to prove I'm an agentic AI expert."
Go in thinking:
"I need to show I can help them decide how far to push agentic AI — and where to stop."
That mindset alone will change how you answer.

One sentence to anchor yourself before the interview
Keep this in your head:
"This interview is about trust, control, and judgement — not intelligence or novelty."
If you answer from that place, you will align naturally with what the Agentic AI Lead is looking for.